Chest X-Ray Segmentation using Deep Neural Network

Parva Gurav (1256570) || Ashwin Emmatty (1255648)
School of Engineering
University of Guelph
Guelph, CA

Abstract :- In this study, we delve into the development of a Deep Neural Network (U-Net) for segmenting Chest X-Ray images. After dataset preparation, we implement the model and conduct training and testing using deep learning techniques. The primary aim of this project is to provide clinicians with a tool for precise segmentation of Chest X-Ray images, aiding in the diagnosis and treatment of thoracic conditions. By automating this process, healthcare professionals can enhance efficiency and accuracy in medical imaging analysis, ultimately improving patient care.

I. Introduction

Lung disorders present a significant challenge in healthcare, encompassing a diverse range of conditions that impair respiratory function. From infections to genetic predispositions, these ailments affect millions worldwide. Among the diagnostic tools available, chest X-ray scans stand out as vital resources, offering quick and non-invasive insights into lung health. They provide crucial information on abnormalities such as tumours and infections, guiding treatment decisions and monitoring disease progression. However, interpreting chest X-rays accurately, particularly in segmenting intricate anatomical structures, poses a notable challenge for conventional techniques. Manual or semi-automated segmentation methods are time-consuming and prone to errors, hindering efficient diagnosis and treatment planning [1].

Deep neural networks (DNNs) have emerged as a transformative solution in medical image analysis, promising unparalleled accuracy and efficiency in segmentation tasks. By leveraging deep learning algorithms, DNNs autonomously learn intricate patterns and features from large datasets, surpassing human capabilities in many instances. In the context of chest X-ray segmentation, DNNs offer the potential to automate and enhance the accuracy of segmentation, thereby improving patient care and radiological diagnosis efficiency. These advanced algorithms streamline the interpretation process, reduce the burden on radiologists, and enable quantitative analysis of imaging biomarkers. Moreover, DNN-based segmentation facilitates objective assessment of disease severity and treatment response, laying the groundwork for precision medicine initiatives.

A. Problem Statement

Accurate interpretation of chest X-ray images is essential for diagnosing various lung disorders effectively. However, conventional segmentation techniques often struggle with the intricate details of anatomical structures, such as the lungs, heart, ribs, and blood vessels. The challenge arises from accurately delineating overlapping structures, which can obscure each other and complicate the segmentation process. Additionally, variations in image quality, including differences in intensity, contrast, and noise levels, further hinder the precise identification of relevant structures.

To address these challenges, this study proposes the adoption of deep neural networks (DNNs), a cutting-edge technology in medical image analysis. DNNs have demonstrated unparalleled capabilities in learning complex patterns and features from vast amounts of data, making them ideally suited for tackling the complexities of chest X-ray segmentation. By leveraging the sophisticated algorithms inherent in DNNs, we aim to automate and significantly enhance the accuracy of segmentation tasks, ultimately improving the diagnostic process and patient care efficiency. This advancement not only reduces the burden on radiologists but also ensures prompt and reliable diagnoses, leading to better patient outcomes. Moreover, by harnessing advanced technology like DNNs, we pave the way for future advancements in medical imaging and diagnostic capabilities [2].

II. Literature Review

Chest radiography, commonly referred to as chest X-ray, remains the cornerstone of diagnostic imaging for pulmonary disorders. While the interpretation of these radiographs traditionally relies on skilled radiologists, the potential for human error underscores the need for computer-aided detection and diagnosis systems. Over the years, advancements in computer vision (CV) methods have paved the way for automated detection of pulmonary abnormalities, with recent developments driven by the emergence of deep learning techniques.

Deep learning, particularly deep convolutional neural networks (CNNs), has revolutionized medical image analysis by leveraging large annotated datasets and powerful computing hardware. These networks have become the preferred choice for tasks such as feature extraction, segmentation, detection, and classification in chest radiography. By automatically learning intricate patterns and features from chest X-ray images, deep CNNs offer the potential for highly accurate and efficient analysis.

A systematic survey conducted by Agrawal and Choudhary focuses on the application of deep learning techniques to chest X-rays for lung segmentation and the detection/classification of pulmonary disorders. This survey aims to provide a comprehensive overview of existing research conducted on publicly available datasets. Notably, the survey includes studies utilizing Generative Adversarial Network (GAN) models for segmentation and classification tasks, highlighting the growing interest in leveraging GANs to address challenges associated with medical data scarcity.

Furthermore, the survey encompasses research conducted prior to the widespread adoption of deep learning models, offering insights into the evolution of techniques and approaches in the field. By consolidating and summarizing existing literature, this survey serves as a valuable resource for researchers and practitioners seeking to understand the current landscape of chest radiography analysis. Additionally, it fills a notable gap in the literature by dedicating specific attention to chest X-rays, a domain that has not been extensively covered in previous surveys.

Overall, the literature reflects a paradigm shift towards the adoption of deep learning techniques in chest radiography analysis. Through continued research and innovation, these approaches hold the promise of enhancing the accuracy, efficiency, and clinical utility of computer-aided diagnosis systems for pulmonary disorders [3].

III. Methodology

A. Dataset

Precise diagnosis depends on accurate segmentation of chest X-ray images, but conventional techniques have trouble handling intricate anatomical structures. This study aims to automate and enhance the accuracy of chest X-ray segmentation by utilizing deep neural networks as an advanced solution. In order to improve patient care and radiological diagnosis efficiency, a reliable and quick segmentation model for lung structures and anomalies in chest X-ray images is being developed [6].

The dataset comprises chest X-ray images and their corresponding masks, categorized into three classes: COVID, Non-COVID, and Normal. The images are pre-processed and normalized to facilitate model training.

B. Objective

The primary objective of this study is to develop a deep neural network (DNN)-based segmentation model for accurately delineating lung structures and anomalies in chest X-ray images. By leveraging the capabilities of deep learning algorithms, we aim to automate the segmentation process and significantly improve accuracy, thereby enhancing patient care and radiological diagnosis efficiency. Specifically, our goals include training the DNN model on a comprehensive dataset of chest X-ray images, optimizing model architecture and parameters, and evaluating model performance through quantitative analysis and validation. Through this research, we seek to contribute to the advancement of computer-aided diagnosis systems in pulmonary imaging and pave the way for more effective and efficient diagnostic workflows in clinical practice.

1. Model Development: Develop a U-Net architecture-based deep learning model suitable for chest X-ray segmentation.
2. Model Training: Model training using a comprehensive dataset of X-ray images of human Chest and corresponding masks.
3. Optimization: Optimize model architecture and parameters to maximize segmentation accuracy and performance.
4. Evaluation: Analyse and validate the trained model's performance quantitatively using a different test dataset.
5. Clinical Application: Assess the model's potential clinical utility and its impact on diagnostic workflows in pulmonary imaging. [4]

C. Architecture

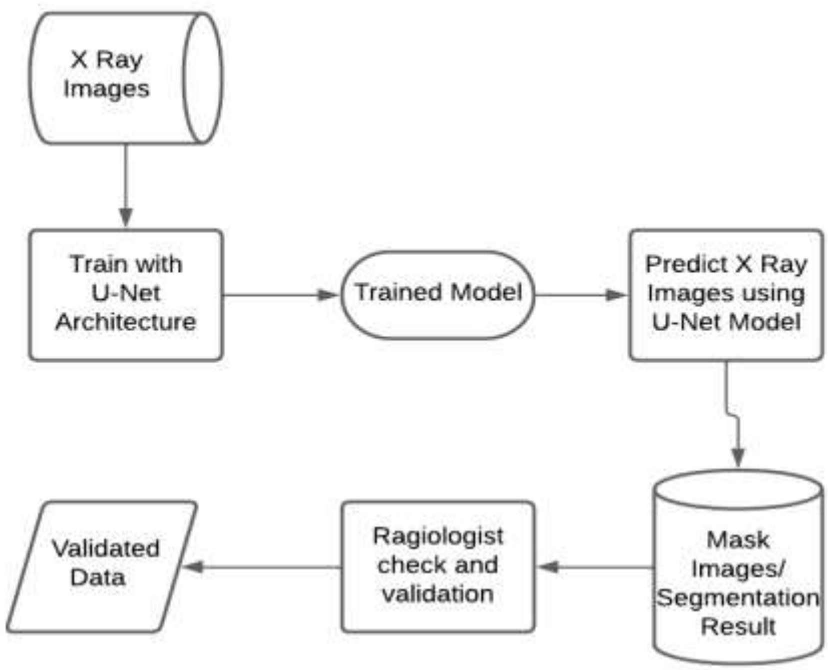

Fig 3.1 – Flow of Project

The proposed U-Net architecture comprises an encoder-decoder structure with skip connections to capture both low-level and high-level features. The encoder consists of convolutional layers followed by max-pooling layers for feature extraction. The decoder employs transposed convolutional layers to upsample the feature maps and reconstruct the segmented image. The final layer utilizes the sigmoid activation function to output pixel-wise segmentation masks [5].

The architecture presented is a U-Net model, a convolutional neural network (CNN) designed for image segmentation tasks, which works great for chest X-ray segmentation. The model consists of an input layer accepting grayscale images of 256x256 pixels, followed by an encoder comprising convolutional layers for feature extraction and max-pooling layers for downsampling. The encoder progressively increases the number of feature maps to capture high-level features. In the decoder, transposed convolutional layers perform upsampling, with skip connections concatenating features from the encoder to maintain spatial information. This enables precise localization and boundary delineation. The final layer utilizes sigmoid activation to generate pixel-wise segmentation masks, representing the probability of belonging to specific classes. With its hierarchical feature learning and skip connections, the U-Net architecture facilitates accurate segmentation of lung structures and anomalies in chest X-ray images, offering significant potential for enhancing medical image analysis and diagnosis.

```
Layer (type)                Output Shape             Param #     Connected to
==================================================================================================
input_2 (InputLayer)        [(None, 256, 256, 1)]    0           []

conv2d_19 (Conv2D)          (None, 256, 256, 32)     320         ['input_2[0][0]']

conv2d_20 (Conv2D)          (None, 256, 256, 32)     9248        ['conv2d_19[0][0]']

max_pooling2d_4 (MaxPoolin  (None, 128, 128, 32)     0           ['conv2d_20[0][0]']
g2D)

conv2d_21 (Conv2D)          (None, 128, 128, 64)     18496       ['max_pooling2d_4[0][0]']

conv2d_22 (Conv2D)          (None, 128, 128, 64)     36928       ['conv2d_21[0][0]']

max_pooling2d_5 (MaxPoolin  (None, 64, 64, 64)       0           ['conv2d_22[0][0]']
g2D)

conv2d_23 (Conv2D)          (None, 64, 64, 128)      73856       ['max_pooling2d_5[0][0]']

conv2d_24 (Conv2D)          (None, 64, 64, 128)      147584      ['conv2d_23[0][0]']

max_pooling2d_6 (MaxPoolin  (None, 32, 32, 128)      0           ['conv2d_24[0][0]']
...
Total params: 7759521 (29.60 MB)
Trainable params: 7759521 (29.60 MB)
Non-trainable params: 0 (0.00 Byte)
```

Fig 3.2 – Architecture of CNN Model

D. ML Approach

The training of the proposed segmentation model employs the Adam optimizer, a popular choice for deep learning tasks, due to its adaptive learning rate properties. The optimizer adjusts the learning rate dynamically during training, allowing for faster convergence and improved performance. In this study, a learning rate of 5e-4 is chosen to balance the trade-off between training speed and convergence stability. To measure the discrepancy between predicted segmentation masks and ground truth masks, the binary cross-entropy loss function is utilized. This loss function quantifies the difference between probability distributions of pixel values in the predicted and ground truth masks, effectively guiding the model towards accurate segmentation.

In addition to the loss function, two key metrics are employed to evaluate the segmentation performance: the Dice coefficient and binary accuracy. The Dice coefficient, also known as the F1 score, measures the spatial overlap between predicted and ground truth masks. It provides a comprehensive assessment of segmentation accuracy by considering both true positive and false positive predictions. The binary accuracy metric, on the other hand, calculates the proportion of correctly classified pixels in the segmentation masks relative to the total number of pixels. Together, these metrics offer insights into the model's ability to accurately delineate lung structures and anomalies in chest X-ray images, facilitating robust evaluation and validation of the segmentation model.

E. Segmentation

During the training phase, the model learns to segment lung structures and anomalies by minimizing the discrepancy between predicted masks and ground truth masks. This iterative optimization process involves adjusting the model's internal parameters, known as weights, based on the gradients of the loss function with respect to these parameters. Utilizing the Adam optimizer with a learning rate of 5e-4, the model iteratively refines its segmentation performance by updating weights to minimize the binary cross-entropy loss function. This loss function quantifies the dissimilarity between predicted segmentation masks and ground truth masks, guiding the model to produce masks that closely resemble the desired segmentations. Through backpropagation, the error signal is propagated backward through the layers of the neural network, enabling the model to learn relevant features and patterns indicative of lung structures and anomalies. By iteratively minimizing the loss function, the model gradually improves its segmentation accuracy and generalization capability, enhancing its ability to accurately segment chest X-ray images.

IV. Results and Discussion

The results of the Chest X-Ray Segmentation project using the U-Net model were promising, demonstrating its effectiveness in accurately delineating anatomical structures. The Dice coefficient, a standard metric for segmentation evaluation, provided valuable insights into the model's accuracy.
The Dice coefficient measures the overlap between predicted and ground truth masks, with values closer to 1 indicating better segmentation performance. By utilizing this metric, we objectively evaluated the model's accuracy (0.9658) in segmenting Chest X-Ray images, guiding improvements for enhanced performance in medical imaging analysis.

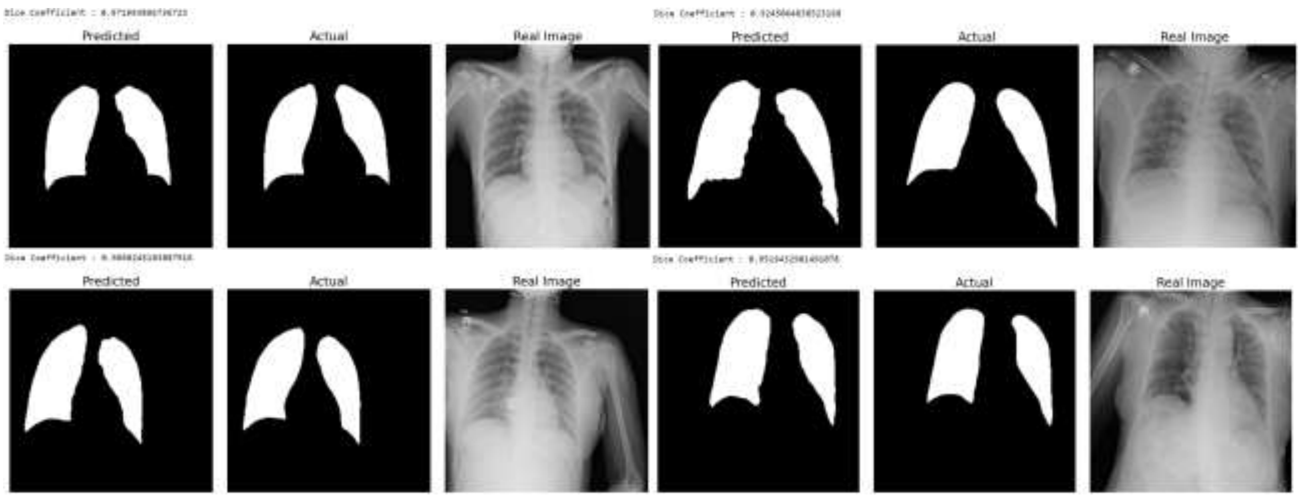

Fig 3.3 – Results of Real Image with Actual and Predicted Image

V. Conclusion

In conclusion, the integration of deep neural networks (DNNs) in chest X-ray segmentation represents a significant leap forward in medical imaging. By harnessing the power of deep learning algorithms, DNNs offer the potential to automate and enhance the accuracy of segmenting intricate anatomical structures, overcoming challenges faced by conventional techniques. The proposed U-Net architecture demonstrates promising results in accurately delineating lung structures and anomalies, paving the way for improved patient care, more efficient radiological diagnosis, and advancements in precision medicine initiatives for pulmonary disorders.

VI. References

[1]. Tarun Agrawal and Prakash Choudhary, "Segmentation and classification on chest radiography: a systematic survey," The Visual Computer, vol. 39, pp. 875-913, Jan. 2022. [Online]. Available: Segmentation and classification on chest radiography: a systematic survey | The Visual Computer (springer.com)

[2]. T. Gabruseva, D. Poplavskiy, and A. A. Kalinin, "Deep Learning for Automatic Pneumonia Detection," in Proceedings of the 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), June 2020, pp. 00183. DOI: 10.1109/CVPRW50498.2020.00183.

[3]. Tarun Agrawal and Prakash Choudhary, "Segmentation and classification on chest radiography: a systematic survey," The Visual Computer, vol. 39, pp. 875-913, Jan. 2022. [Online]. Available: Segmentation and classification on chest radiography: a systematic survey | The Visual Computer (springer.com)

[4]. V. P. H. Vijayan, M. A. Rahiman, and L. Abraham, "Harnessing TernausNet for the Segmentation of Heart from Chest X-ray for the Recognition of Atrial Septal Defect,"

[5]. R. T. Wahyuningrum, I. Yunita, A. Bauravindah, I. A. Siradjuddin, B. D. Satoto, A. K. Sari, and A. D. Sensusiatic, "Chest X-ray Dataset and Ground Truth for Lung Segmentation," Available: https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10570962/

[6]. Rsua, R. (2023, June 19). *RSUA Chest X-Ray Dataset*. Mendeley Data. https://doi.org/10.17632/2jg8vfdmpm.1